

Capstone Project-3

Mobile Price Range Prediction

Navin Kodam

Content :

1. **Defining problem statement**
2. **EDA and feature engineering**
3. **Feature Selection**
4. **Preparing dataset for modeling**
5. **Applying Model**
6. **Model Validation and Selection**
7. **Conclusion**

The Dilemma :

How price range prediction works :

In the competitive mobile phone market companies want to understand sales data of mobile phones and factors which drive the prices. The objective is to find out some relation between features of a mobile phone(eg:- RAM, Internal Memory, etc) and its selling price. In this problem, we do not have to predict the actual price but a price range indicating how high the price is.

Data Summary :

- **Independent variables :**
- **Battery_power** - Total energy a battery can store in one time measured in mAh
- **Blue** - Has bluetooth or not
- **Clock_speed** - speed at which microprocessor executes instructions
- **Dual_sim** - Has dual sim support or not
- **Fc** - Front Camera mega pixels
- **Four_g** - Has 4G or not
- **Int_memory** - Internal Memory in Gigabytes
- **M_dep** - Mobile Depth in cm

Data Summary contd..

- **Mobile_wt** - Weight of mobile phone
- **N_cores** - Number of cores of processor
- **Pc** - Primary Camera mega pixels
- **Px_height** - Pixel Resolution Height
- **Px_width** - Pixel Resolution Width
- **Ram** - Random Access Memory in Mega Bytes
- **Sc_h** - Screen Height of mobile in cm
- **Sc_w** - Screen Width of mobile in cm
- **Talk_time** - longest time that a single battery charge will last when you are

Data Summary contd..

Three_g - Has 3G or not



Touch_screen - Has touch screen or not

Wifi - Has wifi or not

Dependent variables :


Price_range - This is the target variable with value of
0(low cost),
1(mediaum cost),
2(high cost)
and 3(very high cost).

Define Dependent variable :




EDA

Ram vs price_range





EDA contd...

Battery_power vs price_range




EDA contd..

- 3G-4G supported and Non-supported






EDA contd...

Histogram of Front and primary Camera in MegaPixels

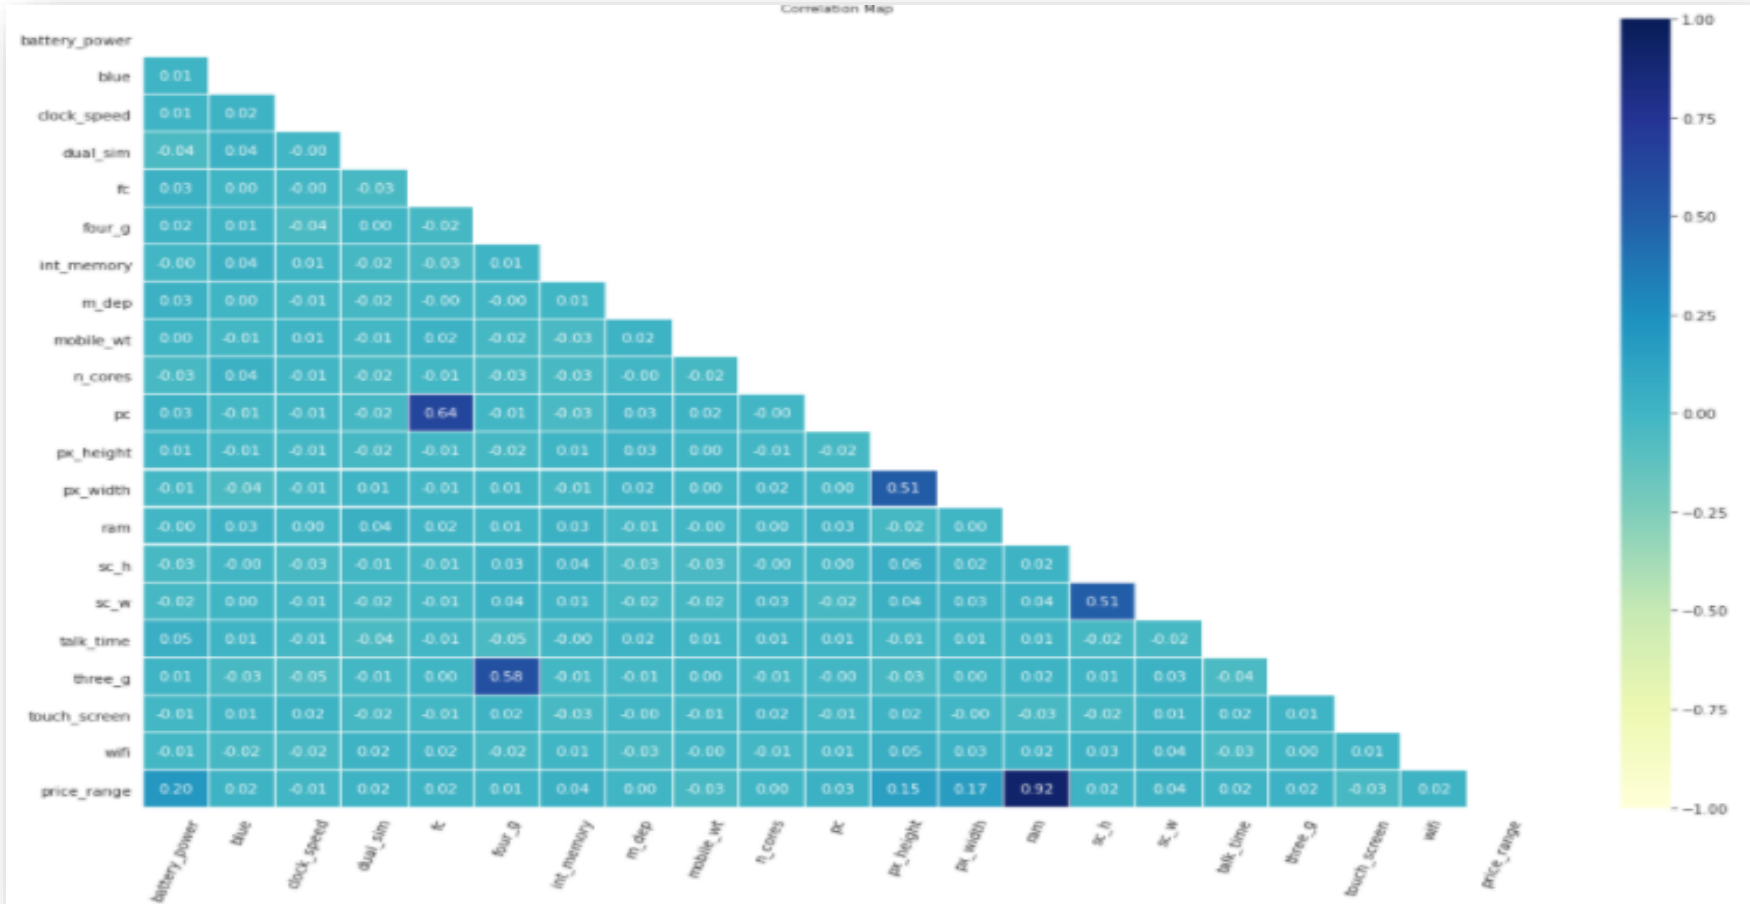


EDA contd...(Plotting negative corr features)

AI

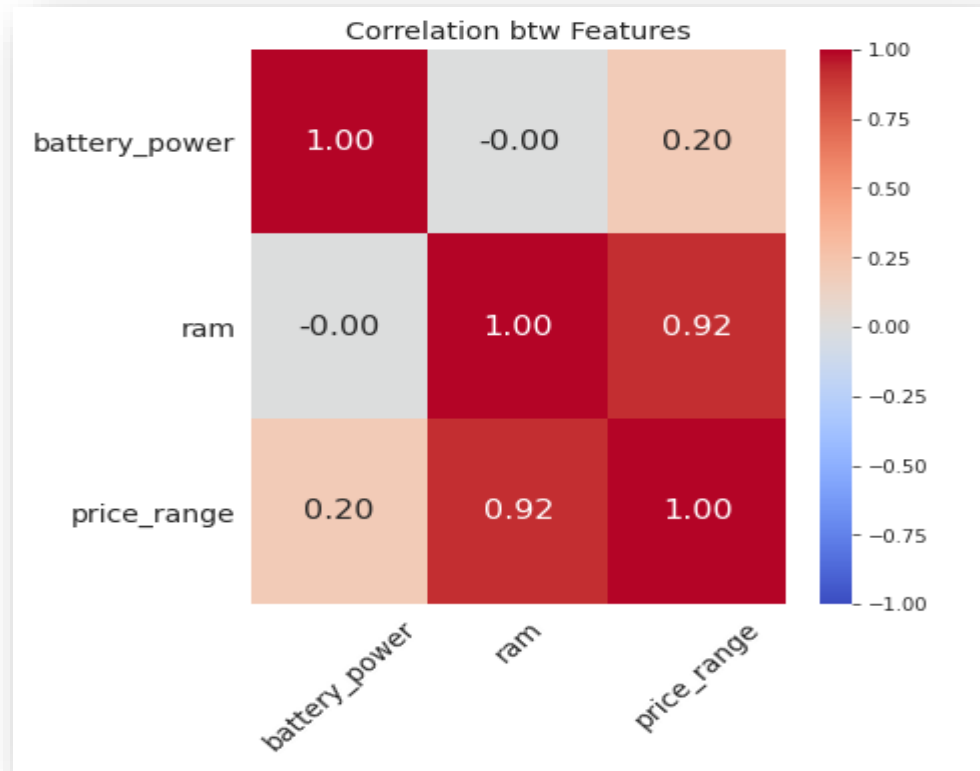


EDA contd...



EDA contd...

Heatmap of maximum Correlation features



Preparing dataset for modeling

**Task : multiclass
classification**

Train set : (1340 , 17)

Test set : (660 , 17)


Response : 0-1-2-3

| battery_power | blue | clock_speed | dual_sim | fc | four_g | int_memory | m_dep | mobile_wt | n_cores | pc | px_height |
|---------------|------|-------------|----------|----|--------|------------|-------|-----------|---------|----|-----------|
| 842 | 0 | 2.2 | 0 | 1 | 0 | 7 | 0.6 | 188 | 2 | 2 | 20 |
| 1021 | 1 | 0.5 | 1 | 0 | 1 | 53 | 0.7 | 136 | 3 | 6 | 905 |
| 563 | 1 | 0.5 | 1 | 2 | 1 | 41 | 0.9 | 145 | 5 | 6 | 1263 |
| 615 | 1 | 2.5 | 0 | 0 | 0 | 10 | 0.8 | 131 | 6 | 9 | 1216 |
| 1821 | 1 | 1.2 | 0 | 13 | 1 | 44 | 0.6 | 141 | 2 | 14 | 1208 |
| 1859 | 0 | 0.5 | 1 | 3 | 0 | 22 | 0.7 | 164 | 1 | 7 | 1004 |
| 1821 | 0 | 1.7 | 0 | 4 | 1 | 10 | 0.8 | 139 | 8 | 10 | 381 |
| 1954 | 0 | 0.5 | 1 | 0 | 0 | 24 | 0.8 | 187 | 4 | 0 | 512 |
| 1445 | 1 | 0.5 | 0 | 0 | 0 | 53 | 0.7 | 174 | 7 | 14 | 386 |
| 509 | 1 | 0.6 | 1 | 2 | 1 | 9 | 0.1 | 93 | 5 | 15 | 1137 |
| 769 | 1 | 2.9 | 1 | 0 | 0 | 9 | 0.1 | 182 | 5 | 1 | 248 |
| 1520 | 1 | 2.2 | 0 | 5 | 1 | 33 | 0.5 | 177 | 8 | 18 | 151 |
| 1815 | 0 | 2.8 | 0 | 2 | 0 | 33 | 0.6 | 159 | 4 | 17 | 607 |

Implementing KNeighbours Classifier

$$\text{TPR} = \text{TP}/(\text{TP}+\text{FN})$$

$$\text{FPR} = \text{FP}/(\text{FP}+\text{TN})$$





Implementing KNeighbours Classifier contd.

Train metrics

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.98 | 0.96 | 0.97 | 228 |
| 1 | 0.93 | 0.96 | 0.94 | 212 |
| 2 | 0.93 | 0.93 | 0.93 | 229 |
| 3 | 0.96 | 0.95 | 0.96 | 228 |
| accuracy | | | 0.95 | 897 |
| macro avg | 0.95 | 0.95 | 0.95 | 897 |
| weighted avg | 0.95 | 0.95 | 0.95 | 897 |

Test metrics

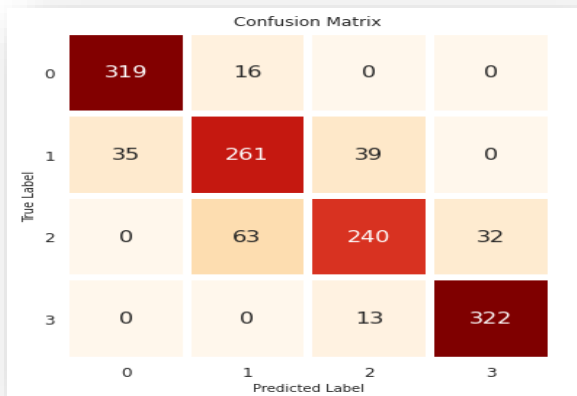
| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.96 | 0.96 | 0.96 | 165 |
| 1 | 0.92 | 0.93 | 0.92 | 165 |
| 2 | 0.88 | 0.90 | 0.89 | 165 |
| 3 | 0.95 | 0.92 | 0.93 | 165 |
| accuracy | | | 0.93 | 660 |
| macro avg | 0.93 | 0.93 | 0.93 | 660 |
| weighted avg | 0.93 | 0.93 | 0.93 | 660 |



Implementing Random Forest Classifier


Train metrics

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.90 | 0.95 | 0.93 | 335 |
| 1 | 0.77 | 0.78 | 0.77 | 335 |
| 2 | 0.82 | 0.72 | 0.77 | 335 |
| 3 | 0.91 | 0.96 | 0.93 | 335 |
| accuracy | | | 0.85 | 1340 |
| macro avg | 0.85 | 0.85 | 0.85 | 1340 |
| weighted avg | 0.85 | 0.85 | 0.85 | 1340 |



Test metrics

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.92 | 0.93 | 0.92 | 165 |
| 1 | 0.71 | 0.76 | 0.74 | 165 |
| 2 | 0.74 | 0.59 | 0.66 | 165 |
| 3 | 0.85 | 0.95 | 0.89 | 165 |
| accuracy | | | 0.81 | 660 |
| macro avg | 0.80 | 0.81 | 0.80 | 660 |
| weighted avg | 0.80 | 0.81 | 0.80 | 660 |



Implementing GradientBoostingClassifier

Train metrics

| Classification Report | | | | | |
|-----------------------|-----------|--------|----------|---------|--|
| | precision | recall | f1-score | support | |
| 0 | 0.91 | 0.94 | 0.93 | 107 | |
| 1 | 0.85 | 0.86 | 0.86 | 123 | |
| 2 | 0.86 | 0.83 | 0.85 | 106 | |
| 3 | 0.94 | 0.93 | 0.94 | 107 | |
| accuracy | | | 0.89 | 443 | |
| macro avg | 0.89 | 0.89 | 0.89 | 443 | |
| weighted avg | 0.89 | 0.89 | 0.89 | 443 | |

Test metrics

| Classification Report | | | | | |
|-----------------------|-----------|--------|----------|---------|--|
| | precision | recall | f1-score | support | |
| 0 | 0.92 | 0.96 | 0.94 | 107 | |
| 1 | 0.88 | 0.85 | 0.86 | 123 | |
| 2 | 0.80 | 0.76 | 0.78 | 106 | |
| 3 | 0.88 | 0.92 | 0.89 | 107 | |
| accuracy | | | 0.87 | 443 | |
| macro avg | 0.87 | 0.87 | 0.87 | 443 | |
| weighted avg | 0.87 | 0.87 | 0.87 | 443 | |

| Confusion Matrix | | | | | |
|------------------|-----|-----------------|----|-----|---|
| True Label | 0 | 1 | 2 | 3 | |
| | 101 | 6 | 0 | 0 | |
| | 10 | 106 | 7 | 0 | |
| | 0 | 12 | 88 | 6 | |
| | 0 | 0 | 7 | 100 | |
| | | 0 | 1 | 2 | 3 |
| | | Predicted Label | | | |

| Confusion Matrix | | | | | |
|------------------|-----|-----------------|----|----|---|
| True Label | 0 | 1 | 2 | 3 | |
| | 103 | 3 | 1 | 0 | |
| | 9 | 104 | 10 | 0 | |
| | 0 | 11 | 81 | 14 | |
| | 0 | 0 | 9 | 98 | |
| | | 0 | 1 | 2 | 3 |
| | | Predicted Label | | | |



Implementing XGBClassifier

Train metrics

| Classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 0.92 | 0.97 | 0.95 | 107 |
| 1 | 0.92 | 0.88 | 0.90 | 123 |
| 2 | 0.87 | 0.85 | 0.86 | 106 |
| 3 | 0.92 | 0.93 | 0.93 | 107 |
| accuracy | | | 0.91 | 443 |
| macro avg | 0.91 | 0.91 | 0.91 | 443 |
| weighted avg | 0.91 | 0.91 | 0.91 | 443 |

Test metrics

| Classification Report | | | | |
|-----------------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 0.91 | 0.94 | 0.93 | 107 |
| 1 | 0.88 | 0.86 | 0.87 | 123 |
| 2 | 0.84 | 0.81 | 0.83 | 106 |
| 3 | 0.90 | 0.92 | 0.91 | 107 |
| accuracy | | | 0.88 | 443 |
| macro avg | 0.88 | 0.88 | 0.88 | 443 |
| weighted avg | 0.88 | 0.88 | 0.88 | 443 |





Implementing Logistic regression

Train metrics

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.90 | 0.88 | 0.89 | 228 |
| 1 | 0.70 | 0.71 | 0.70 | 212 |
| 2 | 0.72 | 0.67 | 0.70 | 229 |
| 3 | 0.85 | 0.90 | 0.87 | 228 |
| accuracy | | | 0.79 | 897 |
| macro avg | 0.79 | 0.79 | 0.79 | 897 |
| weighted avg | 0.79 | 0.79 | 0.79 | 897 |

Test metrics

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.88 | 0.82 | 0.85 | 165 |
| 1 | 0.60 | 0.57 | 0.59 | 165 |
| 2 | 0.58 | 0.63 | 0.60 | 165 |
| 3 | 0.83 | 0.85 | 0.84 | 165 |
| accuracy | | | 0.72 | 660 |
| macro avg | 0.72 | 0.72 | 0.72 | 660 |
| weighted avg | 0.72 | 0.72 | 0.72 | 660 |



Model Validation & Selection contd...

Observations:

1. As seen in the above slides Random forest classifier is not giving great results , GradientBoostingClassifier is bit better than Random forest in recall and precision
2. XGboost classifier is giving the better results than GB but the recall of random forest classifier is somewhat similar
3. KNeighbors is giving the best results among all of the algorithms
4. Logistic regression is giving low results among all of them

Model Validation & Selection contd...


So we had chosen Kneighbors classifier for the prediction and the best hyperparameters obtained are as below

Best hyperparameters :

Train : (algorithm='auto', leaf_size=30, metric='Euclidean', metric_params=None, n_jobs=None, n_neighbors=11, p=2, weights='distance')

Test : (algorithm='auto', leaf_size=30, metric='euclidean', metric_params=None, n_jobs=None, n_neighbors=17, p=2, weights='distance')

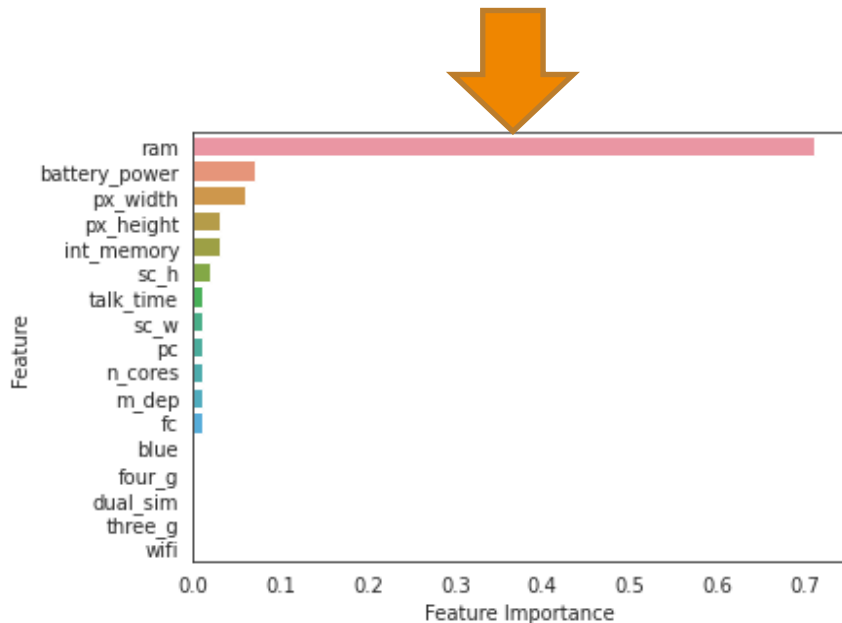
Model Validation & Selection (Hyperparameter tuned)




| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.95 | 0.99 | 0.97 | 228 |
| 1 | 0.90 | 0.92 | 0.91 | 212 |
| 2 | 0.92 | 0.85 | 0.88 | 229 |
| 3 | 0.93 | 0.94 | 0.93 | 228 |
| accuracy | | | 0.92 | 897 |
| macro avg | 0.92 | 0.92 | 0.92 | 897 |
| weighted avg | 0.92 | 0.92 | 0.92 | 897 |

Feature Importance

Random Forest Classifier



XGBoost Classifier



Conclusion

- Ram , Battery_power features were found to be the most relevant features for predicting price range of mobiles and dropping negative correlation features which are clock speed , mobile_wt , touch_screen
- Kneighbors and Xgboost are given best accuracy score 95% test ,93% train and 91% train , 88% test respectively and roc_auc score for kneighbors is 99%
- Tuning the hyperparameters by GridSearchCV on kneighbors but not getting much difference in results but the best parameters n_neighbors for train and test are 11 and 17
- So we conclude that kneighbors classifier is giving the best results for these dataset
- So we can say that in the price range prediction as the ram and battery_power increases the price range will increase for sure

Q & A